

## Assignment 3 (tentative due 24.06)

### Understanding file management.

The goal of this lab is to build a better understanding about how files are treated in UNIX-like OS (we will look at the specific example in Linux). In particular, as you already should know each process maintains a table of open files (file descriptors), later these tables inside processes point to a global table of open files, where at least the current position in the file and the i-node number are maintained. The third table contains a table of inodes, the internal structure of inodes we discussed during the lecture.

To deal with the questions below, create a separate directory, where you copy all given files 1-5.c (the links can be found below near each question).

1.1 Create a directory named INODE. Inside create a file named **hello.txt**:

```
kirillk@OS20:~/OS20/HW3$ cat >> hello.txt
```

Hello world!

!!! Do not forget to press `ctrl-D` to create the end of the file.

To check the **inode number** (in red) run the following:

```
kirillk@OS20:~/OS20/HW3$ ls -li hello.txt
```

```
1310969 -rw-r--r-- 1 kirillk kirillk 6 Jun  9 15:35 hello.txt
```

**I-node number**

When a process opens a file, OS returns a file descriptor that the process uses later to access/change the file.

Copy into the previously created INODE directory (where you have created hello.txt) [1.c](#) and compile it into a.out

- Review 1.c. In general this program opens the previously created file hello.txt containing and prints its process id, file name, and the file descriptor.
- Open two terminals and change directory in both to INODE.
- In one terminal run `./a.out` and copy the printed pid.

```
kirillk@OS20:~/OS20/HW3/INODE$ ./a.out
```

```
pid = 15812
```

```
hello.txt, fd = 3
```

In the second terminal run (do not forget sudo): `sudo lsof -o -p 15812`

```

kirillk@OS20:~$ sudo lsof -o -p 15812
lsof: WARNING: can't stat() fuse.gvfsd-fuse file system /run/user/1000/gvfs
Output information may be incomplete.
COMMAND  PID    USER  FD   TYPE DEVICE  OFFSET      NODE NAME
a.out    15812  kirillk  cwd   DIR    8,1         1181168 /home/kirillk/INODE
a.out    15812  kirillk  rtd   DIR    8,1          2 /
a.out    15812  kirillk  txt   REG    8,1         1181590 /home/kirillk/INODE/a.out
a.out    15812  kirillk  mem   REG    8,1         1971990 /lib/x86_64-linux-gnu/libc-2.27.so
a.out    15812  kirillk  mem   REG    8,1         1971962 /lib/x86_64-linux-gnu/ld-2.27.so
a.out    15812  kirillk   0u    CHR  136,1       0t0      4 /dev/pts/1
a.out    15812  kirillk   1u    CHR  136,1       0t0      4 /dev/pts/1
a.out    15812  kirillk   2u    CHR  136,1       0t0 i-node 4 /dev/pts/1 file name
a.out    15812  kirillk   3u    REG    8,1       0t7 1208013 /home/kirillk/INODE/hello.txt
kirillk@OS20:~$

```

As you can see: 3 means file descriptor 3, **u** - read/write (**r**- if read or **w** if write) t7 means that the current offset in hello.txt is 7 (check lseek in 1.c), also you can see the i-node number 1208013.

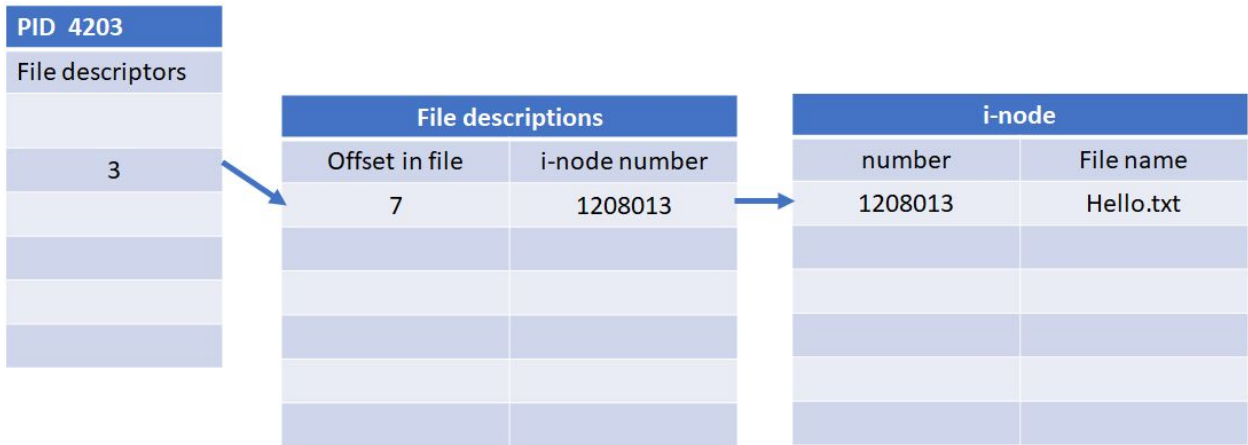


Table 1

As you can see the table of open files contains a file descriptor 3 for hello.txt, this entry points to the global table of open files for all processes containing at least i-node number 1208013 and the current position 7 in the file. Later each such entry points to the third table with i-nodes containing at least i-node number and the file name. For simplicity we use only the local name and not the full name.

In the upcoming questions you will encounter various scenarios. Given the output of the lsof command add its output to the specific places and feel the tables. You will understand when separate entries are created in the first and the second tables (from the left).

**Question 1.1 (25 points):** What happens if we duplicate a file descriptor by dup().

dup() duplicates a file descriptor (see man dup for more details).

Fill Table 2 after running [2.c](#) Review its code before.

PID <your pid>	File descriptions		i-node	
File descriptors	Offset in file	i-node number	number	File name
				hello.txt

Table 2

Put here screenshot of lsof output

**Question 1.2 (25 points):** Fill Table 3 below for the case when the same process opens the same file twice as in 3.c. Review its code. Add the output of lsof to the dedicated place and given the values from lsof fill the table.

PID <your pid>	File descriptions		i-node	
File descriptors	Offset in file	i-node number	number	File name
				hello.txt

Put here screenshot of lsof output

**Question 1.3 (25 points):** Fill Table 4 below if a process opens a file and later fork() as in [4.c](#). Review the code and grab the output of lsof from both processes and based on the values feel the table

PID <your pid>	File descriptions		File descriptions		i-node	
			Offset in file	i-node number	number	File name
						hello.txt

Table 4

Put here screenshot of lsof output for the parent process.

Put here screenshot of lsof output for the child process.

**Question 1.4 (25 points):** Fill Table 5 when two processes as in 5.c open the same file: one run with ./a.out 4 , the other with ./a.out 6. 4 and 6 are the values for lseek in each process. Exit both processes only after you grab lsof from both processes, so use 3 terminals for this purpose.

PID <your pid>
File descriptors

File descriptions	
Offset in file	i-node number

i-node	
number	File name
	hello.txt

PID <your pid>
File descriptors

Table 5

Put here a screenshot of lsof output for the first process with lseek 4

Put here a screenshot of lsof output for the second process with lseek 6